



Financial Due Diligence Kickoff Brief

NVIDIA Corporation (NVDA) —
Semiconductors & Accelerated Computing
January 2026



01

Executive Summary

Financial Snapshot

NVIDIA FY2025 results reflect historic growth driven by Data Center AI demand; margin compression emerging in 9M FY2026

Financial Snapshot		
Metric	Value	Delta
FY2025 Revenue	\$130.5B	+114% YoY
Gross Margin (Peak)	75.0%	↓ to 69.3% in 9M FY2026
Operating Cash Flow	\$64.1B	+128% YoY
Top-4 Cust. Conc.	61%	Q3 FY2026

Key takeaways

- Key takeaways synthesized from the table inputs for quick scan and discussion.
- Use as a working view; confirm all figures during confirmatory diligence.
- Add diligence commentary to contextualize table outcomes and implications.
- Add diligence commentary to contextualize table outcomes and implications.

Key Headlines

Preliminary analysis of public filings and industry data reveals five critical themes that will shape the scope of confirmatory diligence

Revenue Explosion: NVIDIA more than doubled its revenue in FY2025, generating \$130.5B compared to \$60.9B in the prior year. The first nine months of FY2026 have already reached \$147.8B, placing the company on a run-rate approaching \$200B.

Margin Compression: Gross margins peaked at 75.0% in FY2025 but have since compressed to 69.3%, driven by unfavorable shifts in product mix, elevated inventory provisions, and rising input costs across packaging and memory.


Customer Concentration: Four direct customers each accounted for more than 10% of quarterly revenue in Q3 FY2026, creating meaningful counterparty exposure tied to the hyperscaler capital expenditure cycle.

Inventory Build: Inventory nearly doubled from \$10.1B to \$19.8B over just nine months, as the company stages components ahead of the Blackwell platform transition and commits to long-lead supply agreements.

Aggressive Capital Return: The company repurchased \$34.0B of shares in FY2025 and a further \$36.7B in the first nine months of FY2026, signaling management's confidence in the durability of free cash flow generation.

Diligence focus: confirm revenue sustainability, backlog conversion, and pricing discipline.

- Validate margin trajectory, mix shift, and inventory provisioning methodology.
- Assess customer concentration, renewal cadence, and contract terms for key accounts.
- Quantify working capital impacts from lead times, supply commitments, and buffer stock.




Source: NVIDIA 10-K filings

Key Headlines (cont.)

Preliminary analysis of public filings and industry data reveals five critical themes that will shape the scope of confirmatory diligence

- Stress-test downside scenarios tied to hyperscaler capex cycles and platform transitions.



Source: NVIDIA 10-K filings

Company & Transaction Overview

NVIDIA operates as a fabless semiconductor and systems company whose economics are overwhelmingly driven by Data Center accelerated computing


Business Model: NVIDIA designs silicon architectures and full-stack software platforms while relying on TSMC for wafer fabrication, external partners for assembly and test, and SK hynix as its primary supplier of high-bandwidth memory.

- The company reports through two operating segments: Compute & Networking, which generated \$116.2B in FY2025, and Graphics, which contributed \$14.3B.

Market Position: Gartner ranked NVIDIA as the largest semiconductor vendor globally by revenue in 2024, a position underpinned by the CUDA software ecosystem, which has accumulated more than six years of developer tooling and creates substantial switching costs for enterprise customers.

Supply Chain: The company faces structural supply constraints in TSMC's CoWoS advanced packaging and in high-bandwidth memory allocation, both of which limit near-term production capacity.

- TSMC is expanding CoWoS output to approximately 75,000 wafers per month in 2025, roughly doubling 2024 levels, though customer demand continues to outpace available supply.



Source: NVIDIA FY2025 10-K (\$130.5B total)

Valuation Context — Peer Comparison

NVIDIA trades at a significant premium to semiconductor peers on all traditional valuation metrics, reflecting the market's expectations for sustained hypergrowth and platform dominance

Valuation Context — Peer Comparison					
Metric	NVIDIA	AMD	Broadcom	Marvell	Peer Median
Market Cap (\$B)	3,420	202	1,040	98	202
EV / Revenue (NTM)	17.2x	8.4x	15.6x	12.8x	12.8x
EV / EBITDA (NTM)	24.8x	23.1x	25.4x	32.6x	25.4x
P/E (NTM)	30.5x	25.2x	30.8x	42.1x	30.8x
PEG Ratio	0.7x	1.2x	1.5x	2.1x	1.5x
Revenue Growth (NTM)	+48%	+28%	+17%	+24%	+24%
Gross Margin (LTM)	69.3%	49.0%	63.2%	46.8%	49.0%
ROIC	115%	12%	18%	8%	12%

Key takeaways

Valuation premium is clear: EV/Revenue 17.2x vs peer median 12.8x.

- Earnings multiple embeds durability expectations: P/E 30.5x vs peer median 30.8x.
- Premium is underwritten by forward growth assumptions (+48% vs peers ~+24%).

Key diligence: sustainability of gross margin, customer concentration, and supply chain constraints.

Implied Enterprise Value — Sensitivity Analysis

Sensitivity matrix illustrating the range of implied enterprise values at varying Adjusted EBITDA and EV/EBITDA assumptions

Implied Enterprise Value — Sensitivity Analysis					
Adj. EBITDA \ Multiple	20x	24x	28x	32x	36x
Bear (\$70B)	\$1,400B	\$1,680B	\$1,960B	\$2,240B	\$2,520B
Base (\$86B)	\$1,720B	\$2,064B	\$2,408B	\$2,752B	\$3,096B
Bull (\$100B)	\$2,000B	\$2,400B	\$2,800B	\$3,200B	\$3,600B
Upside (\$120B)	\$2,400B	\$2,880B	\$3,360B	\$3,840B	\$4,320B

Key takeaways

- Range of implied enterprise values is highly sensitive to both normalized EBITDA and multiple selection.
- Focus analysis on downside cases (mix/price, margin normalization) and upside scenarios (capacity unlocks).
- Use as a decision framework for valuation guardrails and diligence priorities.
- Add diligence commentary to contextualize table outcomes and implications.

02

Financial Overview

Financial Summary

Revenue more than doubled in FY2025; gross margin peak of 75.0% now compressing as product mix and costs shift

Financial Summary				
\$ in billions	FY2023	FY2024	FY2025	9M FY2026
Revenue	27.0	60.9	130.5	147.8
Gross profit	n/a	44.3	97.9	102.4
Gross margin %	n/a	72.7%	75.0%	69.3%
Operating income	4.2	33.0	81.5	92.6
Operating margin %	15.6%	54.2%	62.5%	62.7%
Net income	4.4	29.8	72.9	77.1
Operating cash flow	5.6	28.1	64.1	n/a
Share repurchases	n/a	n/a	34.0	36.7
Inventory (period-end)	n/a	n/a	10.1	19.8

Key takeaways

- Revenue scale step-change (27.0 → 130.5) with continued momentum into FY2026.
- Margins peaked in FY2025 (75.0%) and moderated in 9M FY2026 (69.3%), requiring mix and provision diligence.
- Cash generation remains strong (operating cash flow 64.1), supporting aggressive capital return.
- Inventory build (10.1 → 19.8) elevates focus on reserves, backlog quality, and platform transitions.

Revenue by End Market


Data Center dominance continues to accelerate, representing 88% of FY2025 revenue, while Gaming and all other segments become increasingly marginal to the business

FY2025 Revenue Composition (\$130.5B)

- The Data Center segment generated \$115.2B in FY2025, representing 88.3% of total revenue. Within Data Center, Compute contributed \$102.2B and Networking added \$13.0B.
- Gaming contributed \$11.35B (8.7%), Professional Visualization accounted for \$1.88B (1.4%), Automotive generated \$1.69B (1.3%), and OEM & Other represented the remaining \$0.39B (0.3%).

Nine Months FY2026 (\$147.8B, annualizing to ~\$197B)


- Through the first nine months of FY2026, Data Center revenue reached \$131.4B, or 88.9% of the total. Networking revenue accelerated to \$20.4B within this period, growing meaningfully faster than the Compute subsegment as hyperscalers deployed increasingly large GPU clusters requiring denser interconnect.



Source: NVIDIA 10-K filings; KPMG analysis

Data Center Revenue Decomposition

Networking revenue is growing faster than Compute, reflecting NVIDIA's platform attach strategy and surging demand for NVLink and InfiniBand interconnect




Source: NVIDIA 10-K/10-Q; FY2024 compute/networking split estimated

- Data Center compute revenue grew from approximately \$34.4B in FY2024 to \$102.2B in FY2025, a 197% increase driven primarily by H100 and B100 platform adoption across hyperscaler and enterprise customers.
- Networking revenue reached \$20.4B in the first nine months of FY2026 alone, implying a run-rate of approximately \$27B. This acceleration reflects the growing importance of high-speed interconnect as customers deploy larger, multi-rack GPU clusters.
- Public filings do not disclose units shipped by platform, average selling price waterfalls, or the revenue mix between discrete boards, integrated systems, and full rack configurations — all of which represent critical data requests for confirmatory diligence.

Data Center Revenue Decomposition (cont.)

Networking revenue is growing faster than Compute, reflecting NVIDIA's platform attach strategy and surging demand for NVLink and InfiniBand interconnect




- Average selling prices have shifted materially as the product portfolio evolves from discrete GPU boards priced in the \$25,000–\$40,000 range to HGX multi-GPU systems at \$200,000–\$400,000 and full DGX SuperPOD rack configurations exceeding \$3M per unit. This ASP migration creates both revenue leverage and margin complexity as the bill of materials scales with system-level integration.

Source: NVIDIA 10-K/10-Q; FY2024 compute/networking split estimated

Data Center Revenue Decomposition (cont.)

Networking revenue is growing faster than Compute, reflecting NVIDIA's platform attach strategy and surging demand for NVLink and InfiniBand interconnect




- Networking revenue per GPU deployed is emerging as a critical efficiency metric for diligence. Each Blackwell-generation cluster requires significantly higher NVLink and InfiniBand interconnect density than prior generations, and NVIDIA's ability to capture this attach revenue represents a durable competitive advantage that custom ASIC alternatives cannot easily replicate.

Source: NVIDIA 10-K/10-Q; FY2024 compute/networking split estimated

Quarterly Revenue Trajectory

Revenue growth rates are decelerating from a peak of +265% year-over-year, but absolute dollar accretion remains unprecedented in semiconductor history




Source: NVIDIA quarterly filings; fiscal year ends January

- NVIDIA reported Q3 FY2026 revenue of \$55.5B, compared to \$18.1B in the year-ago quarter, representing a 206% increase. Sequential growth of 11% suggests the trajectory is beginning to mature as the installed base scales.
- The FY2025 quarterly progression from \$26.0B to \$30.0B to \$35.1B to \$39.3B illustrates a 51% ramp across the fiscal year, with each successive quarter setting a new record for any semiconductor company.
- The nine-month FY2026 run-rate implies approximately \$197B in annualized revenue. Sustaining this trajectory will require continued capital expenditure commitments from the major hyperscalers, which represent the dominant share of Data Center demand.

Quarterly Revenue Trajectory (cont.)

Revenue growth rates are decelerating from a peak of +265% year-over-year, but absolute dollar accretion remains unprecedented in semiconductor history




Source: NVIDIA quarterly filings; fiscal year ends January

- Year-over-year growth rates have decelerated from a peak of approximately 265% in mid-FY2025 to roughly 94% by Q3 FY2026, a pattern consistent with the mathematical base effect of a rapidly scaling business. Sequential growth has also moderated from approximately 15–17% in earlier quarters to 10–11% in Q3 FY2026, suggesting the trajectory is beginning to mature.

Quarterly Revenue Trajectory (cont.)

Revenue growth rates are decelerating from a peak of +265% year-over-year, but absolute dollar accretion remains unprecedented in semiconductor history




Source: NVIDIA quarterly filings; fiscal year ends January

- Management has historically guided conservatively relative to reported results, with actual quarterly revenue exceeding midpoint guidance by an average of \$2–3B over the past six quarters. Backlog visibility is estimated at approximately two to three quarters based on the long lead times associated with CoWoS and HBM supply commitments, providing meaningful near-term demand certainty.

Hyperscaler Capital Expenditure Overlay

NVIDIA's revenue trajectory is fundamentally linked to the capex commitments of its four largest hyperscaler end customers, whose combined spending is approaching \$300B annually




Source: Company filings, FactSet estimates; NVIDIA fiscal years aligned to calendar; KPMG analysis

- The four largest hyperscaler customers — Microsoft, Amazon, Google, and Meta — collectively invested approximately \$185B in capital expenditure during calendar year 2024 and have guided to approximately \$280B for 2025, representing roughly 50% year-over-year growth in infrastructure spending.
- NVIDIA captures an estimated 25–35% of incremental hyperscaler capex as AI-related infrastructure spending, implying that each \$10B increase in combined hyperscaler capex translates to approximately \$2.5–3.5B of incremental NVIDIA revenue. This relationship provides a useful cross-check against the company's bottom-up revenue projections.

Hyperscaler Capital Expenditure Overlay (cont.)

NVIDIA's revenue trajectory is fundamentally linked to the capex commitments of its four largest hyperscaler end customers, whose combined spending is approaching \$300B annually




Source: Company filings, FactSet estimates; NVIDIA fiscal years aligned to calendar; KPMG analysis

- The sustainability of this capex cycle is the single most important assumption underlying NVIDIA's revenue trajectory. Any meaningful pullback in hyperscaler spending — whether driven by ROI concerns, macroeconomic conditions, or a shift toward custom silicon — would have an outsized impact on demand visibility and order backlog.
- On an individual basis, Microsoft is estimated to invest approximately \$80B in calendar 2025, Amazon approximately \$75B, Google approximately \$65B, and Meta approximately \$60B. Each hyperscaler's AI infrastructure investment is increasingly driven by competitive necessity rather than discretionary allocation, which provides a degree of demand resilience but also raises terminal growth rate questions.

Hyperscaler Capital Expenditure Overlay (cont.)

NVIDIA's revenue trajectory is fundamentally linked to the capex commitments of its four largest hyperscaler end customers, whose combined spending is approaching \$300B annually



- The return on invested capital for hyperscaler AI infrastructure remains difficult to quantify at this stage of the deployment cycle. Emerging indicators suggest that inference-driven revenue from AI-powered products and services is growing rapidly but has not yet reached levels sufficient to fully justify the aggregate capital deployed, creating a long-term demand sustainability risk that diligence should stress-test.

Source: Company filings, FactSet estimates; NVIDIA fiscal years aligned to calendar; KPMG analysis

03

Quality of Earnings

Gross Margin Bridge & Sustainability


Gross margin peaked at 75.0% in FY2025 and has compressed by approximately 570 basis points in FY2026, driven by product mix shifts, inventory provisions, and input cost inflation

Margin Drivers: The FY2025 gross margin peak reflected a favorable Data Center revenue mix, premium pricing on supply-constrained GPU platforms, and minimal inventory provision charges during a period of acute demand.

Current Year Compression: The product mix is shifting toward higher-complexity systems and full-rack configurations, which carry a meaningfully higher bill of materials and compress unit margins relative to discrete GPU boards.

- Inventory provisions have increased as the company manages platform transitions from Hopper to Blackwell, while rising costs for HBM memory and CoWoS advanced packaging place additional pressure on input economics.

Peer Benchmark: AMD reported a GAAP gross margin of 49% in FY2024 (53% on a non-GAAP basis). While NVIDIA's current margin regime is exceptional by industry standards, some degree of convergence toward peer levels should be expected as competitive dynamics evolve.



Source: NVIDIA quarterly filings; AMD FY2024 10-K; KPMG analysis

Quality of Earnings — Candidate Adjustments

Preliminary identification of normalization items; magnitude and treatment to be validated during confirmatory diligence

Quality of Earnings — Candidate Adjustments					
Adjustment Category	FY2024	FY2025	9M FY2026	Direction	Priority
SBC — equity awards (RSU/PSU)	2.8	3.7	4.1	Add-back	High
SBC — ESPP & other programs	0.5	0.7	0.7	Add-back	High
Inventory provisions — obsolescence	0.6	0.7	1.6	Normalize	High
Inventory provisions — purchase commits	0.3	0.4	0.8	Normalize	High
Gains/losses on equity investments	(0.2)	(0.4)	(0.3)	Remove	Medium
Gains/losses on debt securities	(0.1)	(0.2)	(0.1)	Remove	Medium
Acquisition amortization (intangibles)	0.1	0.1	0.1	Add-back	Medium
Acquisition integration & deal costs	—	0.1	—	Add-back	Medium
Restructuring / severance	—	0.1	—	Add-back	Low
Lease exit & impairment charges	—	—	0.1	Add-back	Low

Key takeaways

Largest recurring adjustment driver: stock-based compensation (material and growing).

- Inventory-related normalization requires evidence on reserve methodology and platform obsolescence risk.
- Prioritize high-impact items first to converge quickly on normalized earnings.
- Add diligence commentary to contextualize table outcomes and implications.

Quality of Earnings — Candidate Adjustments (cont.)

Preliminary identification of normalization items; magnitude and treatment to be validated during confirmatory diligence

Quality of Earnings — Candidate Adjustments (cont.)					
Adjustment Category	FY2024	FY2025	9M FY2026	Direction	Priority
FX remeasurement gains/losses	(0.1)	0.1	(0.1)	Remove	Low
Indicative adj. to EBITDA	~4.0	~5.2	~6.9		
Memo: Adj. EBITDA margin	58.8%	65.7%	64.2%		

Key takeaways

- Indicative EBITDA adjustment impact is meaningful (FY2025: ~5.2), concentrated in a few items.
- Inventory-related normalization requires evidence on reserve methodology and platform obsolescence risk.
- Prioritize high-impact items first to converge quickly on normalized earnings.
- Add diligence commentary to contextualize table outcomes and implications.

Key observations: SBC represents the largest single adjustment at \$4.8B annualized, and is treated as a recurring operating cost given NVIDIA's dependence on equity compensation to attract and retain engineering talent. Inventory provisions are the fastest-growing adjustment category driven by the Hopper to Blackwell platform transition, and should be normalized across a full product cycle. All figures are preliminary and subject to revision upon receipt of management data during confirmatory diligence.

Quality of Earnings — Adjustment Narratives (1/2)

Detailed rationale for each proposed normalization adjustment, with preliminary magnitude estimates to be validated during confirmatory diligence


Stock-Based Compensation (\$3.3B → \$4.8B)

- NVIDIA's stock-based compensation expense has grown from \$3.3B in FY2024 to \$4.4B in FY2025 and \$4.8B in just the first nine months of FY2026. The acceleration reflects both headcount growth in engineering and AI research functions, and the impact of elevated share prices on the fair value of equity grants. In a normalized earnings context, SBC is treated as a recurring operating cost rather than a non-cash add-back, because the company must continue issuing equity to attract and retain technical talent in a hypercompetitive labor market. The dilutive impact and the share repurchase offset of \$34B in FY2025 should be evaluated together to assess the true economic cost to equity holders.

Inventory Provisions, Net (\$0.9B → \$2.4B)

- Inventory provision charges nearly tripled from \$0.9B in FY2024 to an estimated \$2.4B in the first nine months of FY2026, driven by the Hopper-to-Blackwell platform transition and accelerating obsolescence of older GPU architectures. These provisions represent a recurring feature of NVIDIA's business model given the rapid cadence of product transitions — typically every 18 to 24 months — and should be normalized over a full product cycle rather than treated as one-time items. Diligence should request a detailed roll-forward of the inventory reserve, test the accuracy of prior-period estimates against actual write-offs, and assess whether the current reserve balance adequately reflects the risk of excess inventory in the channel.

Gains and Losses on Strategic Investments (\$0.3B → \$0.6B)



Source: Management report (demo data)

Quality of Earnings — Adjustment Narratives (1/2)


Detailed rationale for each proposed normalization adjustment, with preliminary magnitude estimates to be validated during confirmatory diligence

(cont.)

- NVIDIA maintains a portfolio of strategic equity investments in AI infrastructure, autonomous vehicle, and robotics companies, including positions in both public and private entities. Unrealized gains and losses on this portfolio generated a net gain of approximately \$0.6B in FY2025 and \$0.4B in the first nine months of FY2026. These gains are non-operational and episodic in nature, driven by mark-to-market valuations on public holdings and periodic fair value adjustments on private positions. KPMG recommends removing these items entirely from adjusted EBITDA, as they do not reflect the earnings power of the core semiconductor and platform business.
- Assess pricing power, mix shift, and pass-through of supply-chain costs.
- Confirm customer concentration dynamics and renewal visibility.
- Quantify working capital impact from inventory positioning and lead times.
- Stress-test downside scenarios tied to hyperscaler capex cycles.

Diligence focus: validate sustainability of growth drivers and backlog quality.

- Assess pricing power, mix shift, and pass-through of supply-chain costs.



Source: Management report (demo data)

Quality of Earnings — Adjustment Narratives (2/2)

Additional normalization items and their combined implications for sustainable earnings power


Acquisition-Related Costs (\$0.1B → \$0.2B)

- Acquisition-related costs include amortization of acquired intangible assets, earn-out adjustments, and integration expenses stemming from NVIDIA's acquisitions of Mellanox (2020), Cumulus Networks, and several smaller AI software companies. The annual run-rate of approximately \$0.2B reflects the tail end of purchase price amortization schedules and is declining over time. These costs are properly excluded from adjusted EBITDA as they represent non-cash accounting charges that do not reflect ongoing operational economics. However, diligence should confirm whether any deferred consideration or contingent payments remain outstanding that could represent future cash obligations.

Restructuring and Other One-Time Items (\$0.0B → \$0.1B)

- NVIDIA recorded approximately \$0.1B of restructuring-related charges in FY2025, primarily associated with workforce realignment in the Professional Visualization and Automotive segments and facility consolidation following the shift to hybrid work arrangements. These charges are sporadic, immaterial relative to the company's earnings base, and appropriately treated as non-recurring add-backs. The absence of restructuring charges in FY2024 and 9M FY2026 supports the classification of these items as genuinely one-off in nature.

Net Adjusted Impact — Summary



Source: Management report (demo data)


Quality of Earnings — Adjustment Narratives (2/2)

Additional normalization items and their combined implications for sustainable earnings power (cont.)

- In aggregate, the candidate adjustments increase reported EBITDA by approximately \$4.0B in FY2024, \$5.2B in FY2025, and an annualized \$6.9B based on the first nine months of FY2026. The most material adjustments — stock-based compensation and inventory provisions — are inherently recurring costs of doing business in the semiconductor industry and warrant careful treatment in any buyer's financial model. The net normalization impact as a percentage of reported EBITDA ranges from approximately 5% to 7%, which is modest relative to the underlying earnings base and suggests that reported earnings quality is generally high.
- Assess pricing power, mix shift, and pass-through of supply-chain costs.
- Confirm customer concentration dynamics and renewal visibility.
- Quantify working capital impact from inventory positioning and lead times.
- Stress-test downside scenarios tied to hyperscaler capex cycles.

Diligence focus: validate sustainability of growth drivers and backlog quality.


- Assess pricing power, mix shift, and pass-through of supply-chain costs.



Source: Management report (demo data)

FY2025 Adjusted EBITDA Bridge

Walking from reported net income to adjusted EBITDA produces a normalized earnings figure of approximately \$85.8B, representing a 65.7% adjusted EBITDA margin




Source: NVIDIA FY2025 10-K; KPMG preliminary normalization analysis (\$ in billions)

- The bridge begins with reported FY2025 net income of \$72.9B and adds back \$4.6B of income tax expense, \$0.3B of net interest expense, and \$2.8B of depreciation and amortization to arrive at reported EBITDA of approximately \$80.6B.
- KPMG's preliminary normalization adjustments add \$4.4B of stock-based compensation, \$1.1B of net inventory provisions above normalized levels, and \$0.3B of acquisition-related and restructuring costs, while removing \$0.6B of non-operating investment gains, resulting in net adjustments of approximately \$5.2B.

FY2025 Adjusted EBITDA Bridge (cont.)

Walking from reported net income to adjusted EBITDA produces a normalized earnings figure of approximately \$85.8B, representing a 65.7% adjusted EBITDA margin




Source: NVIDIA FY2025 10-K; KPMG preliminary normalization analysis (\$ in billions)

- The resulting Adjusted EBITDA of approximately \$85.8B implies a 65.7% margin on \$130.5B of revenue. This provides a more meaningful basis for valuation multiples and comparable company analysis than reported net income, which benefits from a structurally low effective tax rate.
- From a sensitivity perspective, a 100 basis point change in gross margin equates to approximately \$1.3B of Adjusted EBITDA impact at current revenue levels. Similarly, every \$10B of incremental revenue at current incremental margins contributes approximately \$7.5B of additional Adjusted EBITDA, illustrating the extraordinary volume leverage embedded in the business model.

FY2025 Adjusted EBITDA Bridge (cont.)

Walking from reported net income to adjusted EBITDA produces a normalized earnings figure of approximately \$85.8B, representing a 65.7% adjusted EBITDA margin




- Relative to semiconductor peers, NVIDIA's Adjusted EBITDA margin of 65.7% is approximately 20–25 percentage points above AMD and 30–35 percentage points above Intel. While some degree of margin convergence is expected as competitive dynamics evolve, the structural advantages of NVIDIA's platform ecosystem suggest sustainable margins well above peer levels.

Source: NVIDIA FY2025 10-K; KPMG preliminary normalization analysis (\$ in billions)

Operating Expense Scalability

R&D and SG&A are growing in absolute terms but declining as a percentage of revenue, demonstrating the significant operating leverage inherent in a scaled platform business




Source: NVIDIA 10-K/10-Q; KPMG analysis

- Research and development expense increased from \$8.7B in FY2024 to \$12.9B in FY2025 and reached \$13.3B in the first nine months of FY2026, driven primarily by headcount-related compensation and investment in compute infrastructure for internal AI workloads.
- Selling, general, and administrative expense grew from \$2.7B to \$3.3B to \$2.6B over the same periods. Despite the absolute increase in total operating costs, operating margins expanded from 54.2% in FY2024 to 62.5% in FY2025, underscoring the degree of incremental margin leverage at this revenue scale.

Operating Expense Scalability (cont.)

R&D and SG&A are growing in absolute terms but declining as a percentage of revenue, demonstrating the significant operating leverage inherent in a scaled platform business




Source: NVIDIA 10-K/10-Q; KPMG analysis

- A key normalization question for diligence is how much of the current R&D spend represents maintenance of existing platforms versus investment in next-generation architectures, and what the sustainable run-rate will be once the Blackwell transition is complete.
- R&D spending is concentrated across three primary categories: headcount-related compensation (estimated at approximately 65% of total R&D), compute infrastructure for internal AI model training and validation (approximately 20%), and silicon tape-out costs, mask sets, and prototype fabrication (approximately 15%). Headcount growth has averaged roughly 15–20% annually over the past two years, predominantly in GPU architecture and CUDA software engineering.

Operating Expense Scalability (cont.)

R&D and SG&A are growing in absolute terms but declining as a percentage of revenue, demonstrating the significant operating leverage inherent in a scaled platform business




- SG&A as a percentage of revenue declined from 10.0% in FY2023 to 2.5% in FY2025, reflecting the extreme operating leverage inherent in a platform business model. KPMG will assess the sustainability of this ratio and evaluate whether certain SG&A functions are under-invested relative to the scale and complexity of the business.

Source: NVIDIA 10-K/10-Q; KPMG analysis

FY2025 Income Statement Build

A revenue-to-net-income conversion rate of 55.9% reflects exceptional operating leverage and a cost structure that scales efficiently with volume




Source: NVIDIA FY2025 10-K; KPMG analysis

- NVIDIA converted \$130.5B of FY2025 revenue into \$97.9B of gross profit at a 75.0% margin, \$81.5B of operating income at a 62.5% margin, and \$72.9B of net income at a 55.9% conversion rate — a profitability profile that is extraordinary by any industry standard.
- Below the operating line, the company benefits from approximately \$3.4B in interest income on its substantial cash and investment balances, as well as a structurally low effective tax rate enabled by intellectual property domiciling.

FY2025 Income Statement Build (cont.)

A revenue-to-net-income conversion rate of 55.9% reflects exceptional operating leverage and a cost structure that scales efficiently with volume




Source: NVIDIA FY2025 10-K; KPMG analysis

- From a sensitivity perspective, each 100 basis points of gross margin movement equates to roughly \$1.3B of EBITDA impact at current revenue levels. Volume leverage, rather than pricing or cost control, is the primary driver of earnings quality.
- NVIDIA's effective tax rate has averaged approximately 12–14% over the past three fiscal years, materially below the U.S. statutory rate of 21%. This favorable rate is driven by the company's intellectual property structure, which routes licensing income through subsidiaries in jurisdictions with preferential tax regimes. Diligence should assess the durability of this structure in light of evolving OECD Pillar Two minimum tax rules.

FY2025 Income Statement Build (cont.)

A revenue-to-net-income conversion rate of 55.9% reflects exceptional operating leverage and a cost structure that scales efficiently with volume



- Capital expenditure intensity remains low at approximately 2–3% of revenue, reflecting the fabless semiconductor model. However, NVIDIA is increasingly investing in internal compute infrastructure for AI model development and validation, and this category of spend is growing faster than revenue. The distinction between maintenance and growth capex will be a key focus area for free cash flow normalization.

Source: NVIDIA FY2025 10-K; KPMG analysis

04

Working Capital

Inventory Position & Supply Chain Risk


Inventory nearly doubled in nine months, rising from \$10.1B to \$19.8B, as the company pre-builds components ahead of the Blackwell platform transition

Inventory Build Drivers: The Blackwell platform ramp requires NVIDIA to stage components months ahead of customer shipment, resulting in a significant increase in work-in-progress and raw material balances.

- Long lead times on CoWoS advanced packaging and HBM memory offtake agreements force the company to make forward purchase commitments well in advance of revenue recognition.
- As NVIDIA shifts toward selling integrated systems and full-rack configurations, the average bill of materials per unit has increased, mechanically driving higher inventory levels across all stages of production.

Reserve Quality Questions: Inventory provision charges have historically been lumpy and are sensitive to the timing of platform transitions. Diligence should test the reserve methodology, trigger thresholds, and accuracy of prior-period estimates.


- Off-balance-sheet purchase obligations with TSMC and memory suppliers may represent material commitments that are not fully visible in the reported inventory balance.



Source: NVIDIA balance sheet; SEC filings

Cash Generation & Capital Allocation

Operating cash flow scaled to \$64.1B in FY2025, and the company is deploying capital aggressively through an accelerating share repurchase program




Source: NVIDIA cash flow statement; SEC filings

- NVIDIA generated \$64.1B of operating cash flow in FY2025, more than doubling the \$28.1B produced in FY2024. This represents an EBITDA-to-cash conversion rate of approximately 98%, reflecting the asset-light nature of the fabless model and favorable working capital dynamics during a period of rapid growth.
- The company repurchased \$34.0B of shares in FY2025 and \$36.7B in the first nine months of FY2026, partially offsetting the dilutive impact of \$4.4B in stock-based compensation expense. The pace of buybacks has meaningfully accelerated alongside cash generation.

Cash Generation & Capital Allocation (cont.)

Operating cash flow scaled to \$64.1B in FY2025, and the company is deploying capital aggressively through an accelerating share repurchase program




Source: NVIDIA cash flow statement; SEC filings

- Free cash flow normalization during diligence should isolate maintenance versus growth capital expenditures, distinguish company-funded from customer-funded infrastructure investments, adjust for stock-based compensation settlement timing, and account for working capital cyclicalities.

Cash Generation & Capital Allocation (cont.)

Operating cash flow scaled to \$64.1B in FY2025, and the company is deploying capital aggressively through an accelerating share repurchase program




Source: NVIDIA cash flow statement; SEC filings

- The company's balance sheet held approximately \$43.2B in cash, cash equivalents, and short-term investments as of October 2025, providing substantial financial flexibility. Total debt stands at approximately \$8.5B, resulting in a net cash position of \$34.7B and a net debt-to-EBITDA ratio that is effectively negative. This fortress balance sheet provides significant capacity for additional share repurchases, strategic acquisitions, or organic investment.

Cash Generation & Capital Allocation (cont.)

Operating cash flow scaled to \$64.1B in FY2025, and the company is deploying capital aggressively through an accelerating share repurchase program



- Dividend payments remain nominal at approximately \$0.5B annually, reflecting management's preference for share repurchases as the primary capital return mechanism. The current share repurchase authorization has approximately \$16B remaining, and KPMG expects the board to authorize additional capacity given the pace of buyback execution and the company's cash generation profile.

Source: NVIDIA cash flow statement; SEC filings

Customer Concentration Analysis

Four direct customers each represent more than 10% of Q3 FY2026 revenue, creating structural concentration risk that is fundamentally tied to the hyperscaler capital expenditure cycle


Concentration Profile: In Q3 FY2026, NVIDIA's four largest direct customers represented 22%, 15%, 13%, and 11% of quarterly revenue respectively, with the top two customers alone accounting for more than one-third of total sales.

- On a nine-month basis, two direct customers individually exceeded the 10% disclosure threshold, representing 21% and 13% of revenue.

Diligence Implications: The direct customers listed in SEC filings are predominantly add-in board partners and original design manufacturers. The ultimate end customers are major hyperscalers, which introduces an additional layer of demand concentration not fully visible in reported figures.

- Even modest changes in commercial terms, payment timing, or order volumes with these top counterparties can move quarterly cash flows by billions of dollars.


Substitution Risk: Several of NVIDIA's largest end customers are actively developing proprietary accelerator alternatives, including Google's TPU, Amazon's Trainium, and Microsoft's Maia, all of which could reduce long-term dependence on NVIDIA hardware.



Source: NVIDIA 10-Q (Q3 FY2026); customer identities not disclosed

Net Working Capital Composition

Inventory has become the dominant working capital driver, with net working capital reaching \$27.0B as of October 2025, reflecting the massive Blackwell prebuild cycle




- Accounts receivable increased to \$17.3B as of October 2025, up from \$10.1B at the January 2025 fiscal year-end. The growth reflects larger individual deal sizes and, in some cases, extended payment terms as the company ships increasingly complex, higher-value system configurations.
- Inventory reached \$19.8B, nearly doubling from \$10.1B over the same period. The buildup is driven by component staging for the Blackwell transition and forward capacity hedging on constrained packaging and memory inputs.

Source: NVIDIA balance sheet; SEC filings

Net Working Capital Composition (cont.)

Inventory has become the dominant working capital driver, with net working capital reaching \$27.0B as of October 2025, reflecting the massive Blackwell prebuild cycle




Source: NVIDIA balance sheet; SEC filings

- Accounts payable grew to \$5.4B from \$3.5B, a more modest increase that reflects NVIDIA's limited leverage relative to its key suppliers. TSMC and SK hynix, as dominant providers of fabrication and memory respectively, generally dictate payment terms rather than negotiate them.
- Deferred revenue increased to approximately \$3.8B as of October 2025, up from \$1.9B at fiscal year-end. The growth reflects the expansion of NVIDIA's enterprise software and support offerings, including AI Enterprise licensing, DGX Cloud subscriptions, and multi-year maintenance contracts. This deferred balance represents a growing source of forward revenue visibility and recurring income.

Net Working Capital Composition (cont.)

Inventory has become the dominant working capital driver, with net working capital reaching \$27.0B as of October 2025, reflecting the massive Blackwell prebuild cycle




- Accrued liabilities reached \$8.2B as of October 2025, encompassing customer rebate and incentive programs, warranty reserves, employee-related accruals, and tax provisions. The customer programs component is the largest and most variable element, and normalization of these accruals will require detailed analysis of historical true-up patterns and program terms during confirmatory diligence.

Source: NVIDIA balance sheet; SEC filings

Working Capital Efficiency — DSO / DIO / DPO

Days inventory outstanding has expanded sharply as the company pre-builds for the Blackwell transition, driving the cash conversion cycle from 34 days to 68 days in nine months




Source: NVIDIA balance sheet; KPMG analysis (days calculated from reported balances and annualized revenue)

- Days sales outstanding increased from approximately 25 days in January 2024 to 43 days as of October 2025, reflecting larger deal sizes with hyperscaler customers and, in some cases, accommodative payment terms offered to secure long-term platform commitments.
- Days inventory outstanding expanded dramatically from 32 days to 49 days over the same period, driven by the Blackwell prebuild cycle and forward component purchases. At current revenue run-rates, each additional day of inventory outstanding equates to roughly \$540M of incremental balance sheet investment.

Working Capital Efficiency — DSO / DIO / DPO

(cont.)
Days inventory outstanding has expanded sharply as the company pre-builds for the Blackwell transition, driving the cash conversion cycle from 34 days to 68 days in nine months




- Days payable outstanding remained relatively stable at 20–24 days, as NVIDIA has limited negotiating leverage with its dominant suppliers — TSMC for fabrication and SK hynix for high-bandwidth memory — both of which operate at capacity and dictate commercial terms. The resulting cash conversion cycle of 68 days represents a meaningful deterioration from the 34-day level at fiscal year-end 2025.

Source: NVIDIA balance sheet; KPMG analysis (days calculated from reported balances and annualized revenue)

Working Capital Efficiency — DSO / DIO / DPO

(cont.)
Days inventory outstanding has expanded sharply as the company pre-builds for the Blackwell transition, driving the cash conversion cycle from 34 days to 68 days in nine months




Source: NVIDIA balance sheet; KPMG analysis (days calculated from reported balances and annualized revenue)

- Relative to semiconductor peers, NVIDIA's cash conversion cycle is now above the industry median of approximately 55 days, a reversal from the below-average levels maintained through FY2025. AMD operates at approximately 80–90 days CCC reflecting its own inventory build for MI300X, while Broadcom maintains a more efficient 35–40 day cycle given its different business model and supply chain structure.

Working Capital Efficiency — DSO / DIO / DPO

(cont.)
Days inventory outstanding has expanded sharply as the company pre-builds for the Blackwell transition, driving the cash conversion cycle from 34 days to 68 days in nine months



Source: NVIDIA balance sheet; KPMG analysis (days calculated from reported balances and annualized revenue)

- Diligence should assess whether the elevated CCC represents a temporary condition tied to the Blackwell platform transition or a structural shift reflecting the company's move toward selling integrated systems. The normalization of net working capital as a percentage of revenue will be a critical input to the free cash flow bridge and, by extension, to enterprise valuation.

05

Risk Assessment


Risk Assessment

KPMG has developed a preliminary, scenario-based risk impact assessment across the key business and financial risk categories relevant to this transaction

Critical Risks: Gross margin compression represents the most pervasive risk: each 100 basis points of decline equates to approximately \$1.3B of annual EBITDA impact at current revenue levels.

- Customer concentration creates significant counterparty exposure, with the single largest customer relationship representing approximately \$29B of annualized revenue at risk.
- The \$19.8B inventory balance carries meaningful obsolescence exposure during platform transitions, particularly if Blackwell adoption timelines shift or demand softens.
- U.S. export controls have already constrained revenue from China and could tighten further, potentially limiting the addressable market for high-performance accelerator products.

Structural Advantages: These risks are partially mitigated by NVIDIA's deep ecosystem advantages, including more than six years of CUDA developer lock-in, incremental operating margins exceeding 80%, and FY2025 operating cash flow generation of \$64.1B.



Source: KPMG analysis — indicative scenario-based impact estimates

U.S. Export Controls — China Revenue Impact


U.S. export restrictions have reduced NVIDIA's addressable market in China by an estimated \$8–12B annually, with further tightening under active consideration by the Bureau of Industry and Security

Regulatory Timeline: The Bureau of Industry and Security first imposed advanced AI chip export restrictions in October 2022, targeting NVIDIA's A100 and H100 GPUs. The rules were expanded in October 2023 to cover the A800 and H800 variants that NVIDIA had developed as compliant alternatives, and further tightened in January 2025 with country-level tiering that restricted shipments to an additional 120 nations.

Revenue Impact: NVIDIA disclosed that China-region revenue represented approximately 17% of Data Center revenue in FY2023, or roughly \$2.5B. Following the initial restrictions, the company developed compliance-grade products, but subsequent rule expansions have progressively closed these workarounds and reduced China Data Center revenue to an estimated \$3.1B in FY2025.

- KPMG estimates the cumulative annual revenue impact of current export controls at \$8–12B, based on the company's prior China market share and the growth trajectory of the Chinese AI infrastructure market.

Forward Risk: Pending regulatory proposals could extend restrictions to inference-class chips and limit the licensing of CUDA software to certain jurisdictions. These measures, if implemented, could further reduce NVIDIA's addressable market by an additional \$3–5B annually and create compliance complexity across the global supply chain.



Source: NVIDIA 10-K/10-Q; BIS rulemaking; KPMG analysis (China revenue estimated)

Competitive Positioning — Margin Comparison (1/3)


NVIDIA's margin structure is exceptional vs semiconductor peers; sustainability is the key diligence question

NVIDIA: NVIDIA achieved a 75.0% gross margin in FY2025, compressing to 69.3% in the current year.

- The margin profile reflects platform pricing power and favorable mix, though the trend is now moving downward.

Diligence focus: confirm revenue sustainability, backlog conversion, and pricing discipline.

- Validate margin trajectory, mix shift, and inventory provisioning methodology.
- Assess customer concentration, renewal cadence, and contract terms for key accounts.
- Quantify working capital impacts from lead times, supply commitments, and buffer stock.
- Stress-test downside scenarios tied to hyperscaler capex cycles and platform transitions.



Competitive Positioning — Margin Comparison (2/3)


NVIDIA's margin structure is exceptional vs semiconductor peers; sustainability is the key diligence question

AMD (Peer): AMD reported a 49% GAAP gross margin in FY2024 (53% non-GAAP), with the MI300X gaining traction in Data Center.

- The margin trend is improving as the higher-margin accelerator mix grows.

Diligence focus: confirm revenue sustainability, backlog conversion, and pricing discipline.

- Validate margin trajectory, mix shift, and inventory provisioning methodology.
- Assess customer concentration, renewal cadence, and contract terms for key accounts.
- Quantify working capital impacts from lead times, supply commitments, and buffer stock.
- Stress-test downside scenarios tied to hyperscaler capex cycles and platform transitions.



Competitive Positioning — Margin Comparison (3/3)


NVIDIA's margin structure is exceptional vs semiconductor peers; sustainability is the key diligence question

Intel (Peer): Intel reported approximately 41% GAAP gross margins in FY2024, reflecting its vertically integrated manufacturing model and different product mix.

- The Gaudi accelerator line is gaining some enterprise traction, but margins remain under structural pressure.

Diligence focus: confirm revenue sustainability, backlog conversion, and pricing discipline.

- Validate margin trajectory, mix shift, and inventory provisioning methodology.
- Assess customer concentration, renewal cadence, and contract terms for key accounts.
- Quantify working capital impacts from lead times, supply commitments, and buffer stock.
- Stress-test downside scenarios tied to hyperscaler capex cycles and platform transitions.



Priority Data Requests

Confirmatory diligence data requests ranked by impact on valuation and risk assessment

Priority Data Requests		
Priority	Data Request	Rationale
High	Monthly units by DC platform/SKU and configuration	Converts revenue to operational reality
High	Customer-level revenue and margin bridge for top customers	Validates concentration risk and pricing
High	Supply commitments: foundry, CoWoS, HBM volumes and pricing	Explains growth ceiling + margin sensitivity
High	Inventory roll-forward, reserve methodology and accuracy	Tests earnings quality and downside protection

Key takeaways

- Top priorities center on unit/ASP bridge, customer concentration, supply commitments, and inventory reserves.
- Goal is to translate reported revenue into operational drivers and risk-adjusted durability.
- Sequence requests to unblock valuation model inputs early, then deepen into root-cause drivers.
- Add diligence commentary to contextualize table outcomes and implications.

Priority Data Requests (cont.)

Confirmatory diligence data requests ranked by impact on valuation and risk assessment

Priority Data Requests (cont.)		
Priority	Data Request	Rationale
High	Backlog/bookings, cancellation behavior, lead-times	Demand visibility and cyclical assessment
Medium	Networking attach data (units per GPU cluster)	Platform expansion and moat validation
Medium	Software/services revenue detail and renewals	Determines durability of platform economics
Medium	Channel inventory, pricing programs, returns/credits	Detects channel stuffing risk

Key takeaways

- Top priorities center on unit/ASP bridge, customer concentration, supply commitments, and inventory reserves.
- Goal is to translate reported revenue into operational drivers and risk-adjusted durability.
- Sequence requests to unblock valuation model inputs early, then deepen into root-cause drivers.
- Add diligence commentary to contextualize table outcomes and implications.

Proposed Diligence Workstreams

KPMG proposes four parallel diligence tracks aligned to the key value drivers, with estimated resource allocation reflecting the relative importance of each workstream

Track 1: Revenue Quality (35%)

- This track will build a bottom-up units-by-configuration-by-ASP model, analyze backlog durability, conduct deep-dive analysis on the top customer relationships, and test channel sell-through patterns against reported shipments.

Track 2: Margin & Cost (25%)


- The margin workstream will construct a detailed gross margin waterfall, decompose bill of materials costs by platform, analyze operating expense scalability, and normalize for stock-based compensation.

Track 3: Working Capital (25%)

- This track will examine the inventory build trajectory, test the adequacy of reserve methodologies, age the accounts receivable portfolio, and develop a normalized free cash flow bridge.

Track 4: Risk & Controls (15%)


- The risk workstream will assess the impact of current and potential export controls, evaluate competitive positioning and substitution threats, map supply chain dependencies, and review related-party transactions.



Source: KPMG FDD standard framework; resource allocation is indicative

AI Accelerator Market Sizing

The data center AI accelerator total addressable market is projected to grow from approximately \$45B in 2023 to over \$320B by 2027; NVIDIA's share remains defensible but competitive intensity is rising




Source: Industry analyst estimates compiled by KPMG; illustrative only

- Industry analysts estimate the AI accelerator total addressable market is growing at a 40–50% compound annual rate through 2027, fueled by the expansion of both training and inference workloads across enterprise, sovereign, and hyperscale customers.
- NVIDIA currently commands an estimated 70–80% share of the training accelerator market. The inference market is more fragmented, however, with custom ASIC solutions from major cloud providers beginning to capture a growing share of inference-specific workloads.

AI Accelerator Market Sizing (cont.)

The data center AI accelerator total addressable market is projected to grow from approximately \$45B in 2023 to over \$320B by 2027; NVIDIA's share remains defensible but competitive intensity is rising




- The competitive landscape includes AMD's MI300X and MI350 product lines, Google's TPU architecture, Amazon's Trainium chips, Microsoft's Maia accelerators, and Broadcom's custom ASIC design partnerships — all of which represent credible long-term alternatives to NVIDIA's GPU-based platform.

Source: Industry analyst estimates compiled by KPMG; illustrative only

AI Accelerator Market Sizing (cont.)

The data center AI accelerator total addressable market is projected to grow from approximately \$45B in 2023 to over \$320B by 2027; NVIDIA's share remains defensible but competitive intensity is rising




- The training-versus-inference split is evolving rapidly, with inference workloads estimated to represent approximately 60–70% of total AI compute demand by 2027, up from roughly 40% in 2024. NVIDIA's competitive moat is strongest in training, where the CUDA ecosystem and multi-GPU scaling capabilities create significant switching costs. The inference market is more fragmented and price-sensitive, creating opportunity for alternative architectures.

Source: Industry analyst estimates compiled by KPMG; illustrative only

AI Accelerator Market Sizing (cont.)

The data center AI accelerator total addressable market is projected to grow from approximately \$45B in 2023 to over \$320B by 2027; NVIDIA's share remains defensible but competitive intensity is rising



- Sovereign AI represents an emerging demand vector, with governments in the Middle East, Southeast Asia, and Europe investing in domestic AI compute infrastructure for national security, healthcare, and public administration applications. This market segment is estimated at \$15–25B by 2027 and is less price-sensitive than commercial hyperscale, potentially supporting premium margins for NVIDIA's enterprise-grade platforms.

Source: Industry analyst estimates compiled by KPMG; illustrative only

Proposed Diligence Timeline

Six-week confirmatory FDD engagement aligned to workstream priorities and data availability

Weeks 1–2: Kickoff & Data

- Execute engagement letter and NDA
- Submit priority data request list
- Conduct management kickoff meeting
- Begin public filing analysis and industry benchmarking
- Set up virtual data room access and review protocols

Weeks 3–4: Core Analysis

- Build bottom-up revenue model by platform and customer
- Construct gross margin bridge and BOM cost decomposition
- Complete inventory roll-forward and reserve adequacy testing
- Analyze AR aging and cash conversion cycle dynamics
- Conduct management interviews on key diligence topics

Weeks 5–6: Synthesis

- Develop normalized EBITDA and free cash flow bridges
- Complete working capital analysis and seasonal adjustments
- Finalize export control and regulatory risk assessment
- Prepare preliminary findings presentation for deal team
- Draft comprehensive FDD report with technical appendices

Post-Close Support

- Support purchase price allocation and valuation
- Assist with integration planning for financial systems
- Provide 100-day plan financial monitoring framework
- Ad-hoc support on identified risk mitigation items
- Ongoing advisory on regulatory compliance matters



The contacts at KPMG in connection with this report are:

Rebecca Placeholder

Job Description

T: +1 416 000 0001

E: rebecca@kpmg.ca

Name Placeholder

Job Description

T: +1 416 000 0002

E: name@kpmg.ca

Name Placeholder

Job Description

T: +1 416 000 0003

E: name2@kpmg.ca



© 2026 KPMG LLP, an Ontario limited liability partnership and a member firm of the KPMG global organization of independent member firms affiliated with KPMG International Limited is a private English company limited by guarantee and does not provide services to clients. For more detail about our structure please visit kpmg.com/governance.

The KPMG name and logo are trademarks used under license by the independent member firms of the KPMG global organization.

www.kpmg.com

Document Classification: KPMG Confidential


Appendix — Illustrative Chart (A)

Synthetic data for demo only

- This chart is illustrative and does not reflect actual results.
- Included to expand chart coverage in the migrated deck.

Diligence focus: confirm revenue sustainability, backlog conversion, and pricing discipline.


- Validate margin trajectory, mix shift, and inventory provisioning methodology.
- Assess customer concentration, renewal cadence, and contract terms for key accounts.
- Quantify working capital impacts from lead times, supply commitments, and buffer stock.
- Stress-test downside scenarios tied to hyperscaler capex cycles and platform transitions.



Demo data

Appendix — Illustrative Chart (B)

Synthetic data for demo only



Demo data

- Line chart example using the V2 chart renderer.
- Diligence focus:** confirm revenue sustainability, backlog conversion, and pricing discipline.
- Validate margin trajectory, mix shift, and inventory provisioning methodology.
 - Assess customer concentration, renewal cadence, and contract terms for key accounts.
 - Quantify working capital impacts from lead times, supply commitments, and buffer stock.
 - Stress-test downside scenarios tied to hyperscaler capex cycles and platform transitions.